

Lab5_Multithreading

20307130350 信息安全 陈丹纯

Uthread: switching between threads

实验内容

设计用户级线程系统上下文切换机制并实现。

实验思路

通过xv6 book对进程调度及其上下文的机制的学习，该实验可以仿照进程的上下文切换，给struct thread也增加一个用户级线程的上下文 context，然后在 uthread_switch.s 中实现store old and restore new.

实验代码

1. 在 user/uthread.c 中创建一个 struct context 并在thread中添加一个context字段。(直接 copy kernel/proc.h 的context结构)

```

13 //add for lab
14 struct context {
15     uint64 ra;
16     uint64 sp;
17
18     // callee-saved
19     uint64 s0;
20     uint64 s1;
21     uint64 s2;
22     uint64 s3;
23     uint64 s4;
24     uint64 s5;
25     uint64 s6;
26     uint64 s7;
27     uint64 s8;
28     uint64 s9;
29     uint64 s10;
30     uint64 s11;
31 };
32
33
34 struct thread {
35     char    stack[STACK_SIZE]; /* the thread's stack */
36     int     state;             /* FREE, RUNNING, RUNNABLE */
37     struct  context context;
38 };

```

2. 在 `uthread_switch.S` 中保存和恢复上下嗯 (仿照 `kernel/switch.S`)

```
7
8     .globl thread_switch
9 thread_switch:
10     /* YOUR CODE HERE */
11
12     sd ra, 0(a0)
13     sd sp, 8(a0)
14     sd s0, 16(a0)
15     sd s1, 24(a0)
16     sd s2, 32(a0)
17     sd s3, 40(a0)
18     sd s4, 48(a0)
19     sd s5, 56(a0)
20     sd s6, 64(a0)
21     sd s7, 72(a0)
22     sd s8, 80(a0)
23     sd s9, 88(a0)
24     sd s10, 96(a0)
25     sd s11, 104(a0)
26
27     ld ra, 0(a1)
28     ld sp, 8(a1)
29     ld s0, 16(a1)
30     ld s1, 24(a1)
31     ld s2, 32(a1)
32     ld s3, 40(a1)
33     ld s4, 48(a1)
34     ld s5, 56(a1)
35     ld s6, 64(a1)
36     ld s7, 72(a1)
37     ld s8, 80(a1)
38     ld s9, 88(a1)
39     ld s10, 96(a1)
40     ld s11, 104(a1)
41
```

3. 在 `thread_create` 中初始化 `ra` 和 `sp`。

```

1 void
2 thread_create(void (*func)())
3 {
4     struct thread *t;
5
6     for (t = all_thread; t < all_thread + MAX_THREAD; t++) {
7         if (t->state == FREE) break;
8     }
9     t->state = RUNNABLE;
10    // YOUR CODE HERE
11    t->context.ra = (uint64)func;
12    t->context.sp = (uint64)&t->stack[STACK_SIZE-1];
13 }
14

```

4. 在 thread_schedule 中切换上下文。

```

11 }
12
13 if (next_thread == 0) {
14     printf("thread_schedule: no runnable threads\n");
15     exit(-1);
16 }
17
18 if (current_thread != next_thread) {          /* switch threads? */
19     next_thread->state = RUNNING;
20     t = current_thread;
21     current_thread = next_thread;
22     /* YOUR CODE HERE
23      * Invoke thread_switch to switch from t to next_thread:
24      * thread_switch(??, ??);
25      */
26     thread_switch((uint64)&t->context, (uint64)&current_thread->context);
27 } else
28     next_thread = 0;
29 }
30

```

实验结果

```
qemu-system-riscv64 -machine virt -bios none -kernel kernel/kernel -m 128M -smp 3 -nographic -drive file=fs.img,if=none,format=raw,id=x0 -device virtio-blk-device,drive=x0,bus=virtio-mmio-bus.0
```

```
xv6 kernel is booting
```

```
hart 2 starting
```

```
hart 1 starting
```

```
init: starting sh
```

```
$ uthread
```

```
thread_a started
```

```
thread_b started
```

```
thread_c started
```

```
thread_c 0
```

```
thread_a 0
```

```
thread_b 0
```

```
thread_c 1
```

```
thread_a 1
```

```
thread_b 1
```

```
thread_c 2
```

```
thread_a 2
```

```
thread_b 2
```

```
thread_c 3
```

```
thread_a 3
```

```
thread_b 3
```

```
thread_c 4
```

```
thread_a 4
```

```
thread_b 4
```

```
thread_c 5
```

```
thread_a 5
```

```
thread_b 5
```

```
thread_c 6
```

```
thread_a 6
```

```
thread_b 6
```

```
thread_c 7
```

```

thread_c 89
thread_a 89
thread_b 89
thread_c 90
thread_a 90
thread_b 90
thread_c 91
thread_a 91
thread_b 91
thread_c 92
thread_a 92
thread_b 92
thread_c 93
thread_a 93
thread_b 93
thread_c 94
thread_a 94
thread_b 94
thread_c 95
thread_a 95
thread_b 95
thread_c 96
thread_a 96
thread_b 96
thread_c 97
thread_a 97
thread_b 97
thread_c 98
thread_a 98
thread_b 98
thread_c 99
thread_a 99
thread_b 99
thread_c: exit after 100
thread_a: exit after 100
thread_b: exit after 100
thread_schedule: no runnable threads
$

```

```

make[1]: Leaving directory '/home/olegdn/xv6-labs-2021
== Test uthread ==
$ make qemu-gdb
uthread: OK (5.3s)
(Old xv6.out.uthread failure log removed)
Test assumes thread.txt assumes thread.txt: FAT

```

实验问题

thread_switch只需要保存/恢复被调用者保存的寄存器。为什么？

caller-saved寄存器由调用的C代码保存在堆栈上，switch保存了ra寄存器和sp，可以从新的上下文中恢复寄存器，新的上下文中保存着前一次switch所保存的寄存器值。当switch返回时，它返回到被恢复的ra寄存器所指向的指令，也就是新线程之前调用switch的指令。此外，sp会指向新线程的栈。

实验反思

通过本实验，深入了解了context的切换机制。

Using threads

实验内容

分析

1. There is something wrong in `notxv6/ph.c`, thus making the code do not function as expected when working with multi-threads.
2. Check out the `ph.c` to find out the reasons, especially in `put()` and `insert()`. Write them in the `answers-thread.txt`.
3. Modify the `ph.c` especially `put()` and `get()` to lock and unlock the **critical sections**. Make sure the modified can pass the `ph_safe` by running `make grade` in terminal.
4. Try more times to optimize the performances to pass the `ph_fast`.

原文

more details in [Lab: Multithreading\(mit.edu\)](https://lab.multithreading.mt.edu/).

实验分析

实验代码和answers-thread.txt已上传git。

多线程出错的原因

因为table是global variable，所以当两个及以上多个线程正好在同一个table中put()时，以两个线程为例：

1. **error:** `insert()` 中的插入链表操作可能会导致覆盖。eg.当线程0运行完 `e->next=n` 之后发生 context switch，线程1也执行 `put e->next=n` 之后并继续执行 `*p=e`；再切换为线程0执行 `*p=e`，先执行 `*p=e` 的线程会丢失新添加的key。
2. **warning:** `put()`中执行 `e->value=value` 来更新存在的key时会导致覆盖，导致不同的线程丢失率不同。比如两个线程都要修改同一个 entry，但由于调度先后执行 `e->value=value`，则先执行该语句的会丢失更新的key，另一线程得以保留。因此两个线程对key的丢失不同。
3. **error:** 两个线程一个在执行 `put()` 时恰好执行到 `insert(key,value,&table,table[i])` 时，切换到另一个线程执行 `insert()` 中的 `*p=e*`，如果针对同一个table，再切换回来会导致 `put()` 中的 `insert(args)` 中的 `table` 变化，导致key丢失。

```
78 static void *
79 put_thread(void *xa)
80 {
81     int n = (int) (long) xa; // thread number
82     int b = NKEYS/nthread;
83
84     for (int i = 0; i < b; i++) {
85         put(keys[b*n + i], n);
86     }
87
88     return NULL;
89 }
```



```
31 insert(int key, int value, struct entry **p, struct entry *n)
32 {
33     struct entry *e = malloc(sizeof(struct entry));
34     e->key = key;
35     e->value = value;
36     e->next = n;
37     *p = e;
38 }
39
40 static
41 void put(int key, int value)
42 {
43
44     int i = key % NBUCKET;
45
46     pthread_mutex_lock(lock+i);
47     // is the key already present?
48     struct entry *e = 0;
49     for (e = table[i]; e != 0; e = e->next) {
50         if (e->key == key)
51             break;
52     }
53     if(e){
54         // update the existing key.
55         e->value = value;
56     } else {
57         // the new is new.
58         insert(key, value, &table[i], table[i]);
59     }
60     pthread_mutex_unlock(lock+i);
61
62 }
63
```

Lock方案

根据上面多线程出错原因分析，是对全局变量 `table` 操作时需要互斥。要保证0 missing，只需解决原因1和原因3即 `insert()` 中对entry节点的覆盖对其他线程 `insert()` 或者 `put()` 中 `insert()`。因此有多种加锁方案。

1. 对hash表的每个table的 `insert` (smallest critical section) 都加互斥锁

1. 因为要对于每个table都加锁，定义全局变量 `lock[NBUCKET]`;

```
pthread_mutex_t lock[NBUCKET];
```

2. 在 `main()` 中初始化锁。

```
for(int i=0;i<NBUCKET;i++){  
    pthread_mutex_init(&lock[i],NULL);  
}
```

3. the smallest critical section为:

```
insert(key, value, &table[i], table[i]);  
  
} else {  
    // the new is new.  
    pthread_mutex_lock(&lock[i]);  
    insert(key, value, &table[i], table[i]);  
    pthread_mutex_unlock(&lock[i]);  
}
```

对应的实验结果:

```
utegan@ubuntu:~/xv6-labs-2021$  
utegan@ubuntu:~/xv6-labs-2021$ make ph  
gcc -o ph -g -O2 -DSOL_THREAD -DLAB_THREAD notxv6/ph.c -pthread  
utegan@ubuntu:~/xv6-labs-2021$ ./ph 2  
100000 puts, 3.247 seconds, 30796 puts/second  
0: 0 keys missing  
1: 0 keys missing  
200000 gets, 6.476 seconds, 30885 gets/second  
utegan@ubuntu:~/xv6-labs-2021$ ./ph 4  
100000 puts, 3.985 seconds, 25095 puts/second  
1: 0 keys missing  
2: 0 keys missing  
3: 0 keys missing  
0: 0 keys missing  
400000 gets, 20.798 seconds, 19233 gets/second  
utegan@ubuntu:~/xv6-labs-2021$ make grade  
make clean  
  
make[1]: Leaving directory '/home/utegan/xv6-labs-2021'  
ph_safe: OK (9.7s)  
== Test ph_fast == make[1]: Entering directory '/home/utegan/xv6-labs-2021'  
make[1]: 'ph' is up to date.  
make[1]: Leaving directory '/home/utegan/xv6-labs-2021'  
ph_fast: OK (24.6s)  
== Test barrier == make[1]: Entering directory '/home/utegan/xv6-labs-2021'  
gcc -o barrier -g -O2 -DSOL_THREAD -DLAB_THREAD notxv6/barrier.c -pthr
```

2. 直接对所有table的smallest critical section加锁。

因为无需对每个table分别加锁，所有的table都加同样的锁,只需定义并初始化一个lock。

```
// the new cs new.  
pthread_mutex_lock(lock);  
insert(key, value, &table[i], table[i]);  
pthread_mutex_unlock(lock);  
}
```

结果:

```
make[1]: Leaving directory '/home/utegan/xv6-labs-2021'  
ph_safe: OK (13.2s)  
== Test ph_fast == make[1]: Entering directory '/home/utegan/xv6-labs-2021'  
make[1]: 'ph' is up to date.  
make[1]: Leaving directory '/home/utegan/xv6-labs-2021'  
ph_fast: OK (26.7s)  
elp == Test barrier == make[1]: Entering directory '/home/utegan/xv6-labs-2021'  
gcc -o barrier -g -O2 -DSOL_THREAD -DLAB_THREAD notxv6/barrier.c -pthread  
make[1]: Leaving directory '/home/utegan/xv6-labs-2021'
```

3. 直接对put()加锁。

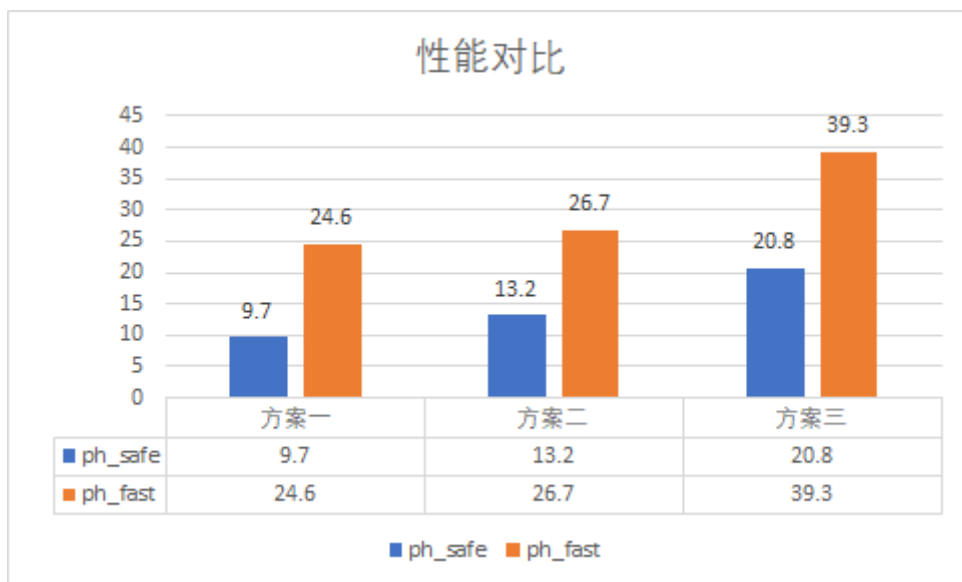
```
for (int i = 0; i < b; i++) {  
    pthread_mutex_lock(lock);  
    put(keys[b*n + i], n);  
    pthread_mutex_unlock(lock);  
}
```

结果:

```
ph_safe: OK (20.8s)  
== Test ph_fast == make[1]: Entering directory '/home/utegan/xv6-labs-2021'  
make[1]: 'ph' is up to date.  
make[1]: Leaving directory '/home/utegan/xv6-labs-2021'  
ph_fast: FAIL (39.3s)  
    Parallel put() speedup is less than 1.25x  
== Test barrier == make[1]: Entering directory '/home/utegan/xv6-labs-2021'  
gcc -o barrier -g -O2 -DSOL_THREAD -DLAB_THREAD notxv6/barrier.c -pthread
```

实验结果分析

不同方案性能分析



原因分析

1. 显然方案三不是最小critical section，其 `ph_fast` 也没有达标。
2. 下面重点分析方案一和方案二：
 1. 两者都是对smallest critical section加锁。
 2. 但是方案一的性能比方案二更高的原因在于方案一**分别对每个table设置了一个特定的锁**，因此当且仅当不同线程对同一个 `table`（即 `i` 相同时）执行 `insert` 时，才会互斥。
 3. 方案二是对所有的 `table` **不加区分地设置了同一个锁**，因此当不同线程执行 `insert` 的时候都会发生互斥，和 `i` 无关，频率更高，对性能的影响自然更大。

Barrier

实验内容

概述

1. read `notxv6/barrier.c`.
2. implement `barrier()` to achieve the desired barrier which is that each thread blocks in `barrier()` until all `nthreads` of them have called `barrier()`.
3. useful primitives

```
pthread_cond_wait(&cond, &mutex);  
pthread_cond_broadcast(&cond);
```

4. skip the assert triggers and pass make grade's barrier test.

原文

more details in [Lab: Multithreading \(mit.edu\)](https://www.mt.edu/~lab/multithreading/).

实验分析

实验代码上传git。

代码报错原因

Each thread executes a loop. In each loop iteration a thread calls `barrier()` and then sleeps for a random number of microseconds. **The assert triggers, because one thread leaves the barrier before the other thread has reached the barrier.**

```
46 static void *
47 thread(void *xa)
48 {
49     long n = (long) xa;
50     long delay;
51     int i;
52
53     for (i = 0; i < 20000; i++) {
54         int t = bstate.round;
55         assert (i == t);
56         barrier();
57         usleep(random() % 100);
58     }
59
60     return 0;
61 }
```

Solution

therefore, to avoid the `assert(i==t)` triggering, the `barrier()` is supposed to make each thread with the same `bstate.round` when running `thread()`'s `block` until all threads get the same `bstate.round`, wake up all threads to go on the next loop.

Algorithm

To realize this goal, in the `barrier()`, increase the `nthread` to account the number of threads in the current round. When `bstate.nthread` is less than the number of all created threads, call `pthread_mutex_wait()`; When equal, increase the `bstate.round` by one and `bstate.nthread` to 0, call `pthread_mutex_broadcast()`;

Note: Don't forget the lock and unlock.

Implement(code)

```
25 static void
26 barrier()
27 {
28     // YOUR CODE HERE
29     //
30     pthread_mutex_lock(&bstate.barrier_mutex);
31     bstate.nthread++;
32     if(bstate.nthread==nthread){
33         bstate.nthread=0;
34         bstate.round++;
35         pthread_cond_broadcast(&bstate.barrier_cond);
36     }else{
37         pthread_cond_wait(&bstate.barrier_cond,&bstate.barrier_mutex);
38     }
39     pthread_mutex_unlock(&bstate.barrier_mutex);
40     // Block until all threads have called barrier() and
41     // then increment bstate.round.
42     //
43
44 }
```

实验结果

```
== Test barrier == make[1]: Entering directory '/home/utegan/xv6-labs-2021'
gcc -o barrier -g -O2 -DSOL_THREAD -DLAB_THREAD notxv6/barrier.c -pthread
make[1]: Leaving directory '/home/utegan/xv6-labs-2021'
barrier: OK (11.8s)
== Test time ==
time: FAIL
    Cannot read time.txt
Score: 24/60
make: *** [Makefile:336: grade] Error 1
utegan@ubuntu:~/xv6-labs-2021$ ./barrier 1
OK; passed
utegan@ubuntu:~/xv6-labs-2021$ ./barrier 2
OK; passed
utegan@ubuntu:~/xv6-labs-2021$ ./barrier 4
OK; passed
utegan@ubuntu:~/xv6-labs-2021$
```

